

Chapter 5

Input/Output

The I/O Subsystem

- The largest, most complex subsystem in OS
- Most lines of code
- Highest rate of code changes
- Where OS engineers most likely to work
- Difficult to test thoroughly
- **Make-or-break issue for any system**
 - » Big impact on performance and perception
 - » Bigger impact on acceptability in market

I/O Devices

- **Types of I/O devices**

- **Block devices**

- » Stores information in fixed-size blocks, each one with its own address
 - » Read or write each block independently
 - » Common block size: 512 bytes to 32,768 bytes
 - » E.g. disks

- **Character devices**

- » Delivers or accepts a stream of characters, without regard to any block structure
 - » Not addressable, no seek operation
 - » E.g. printers, network interfaces, mice

- **Other devices**

- » Clocks, memory-mapped screens

Principles of I/O Hardware

Device	Data rate
Keyboard	10 bytes/sec
Mouse	100 bytes/sec
56K modem	7 KB/sec
Telephone channel	8 KB/sec
Dual ISDN lines	16 KB/sec
Laser printer	100 KB/sec
Scanner	400 KB/sec
Classic Ethernet	1.25 MB/sec
USB (Universal Serial Bus)	1.5 MB/sec
Digital camcorder	4 MB/sec
IDE disk	5 MB/sec
40x CD-ROM	6 MB/sec
Fast Ethernet	12.5 MB/sec
ISA bus	16.7 MB/sec
EIDE (ATA-2) disk	16.7 MB/sec
FireWire (IEEE 1394)	50 MB/sec
XGA Monitor	60 MB/sec
SONET OC-12 network	78 MB/sec
SCSI Ultra 2 disk	80 MB/sec
Gigabit Ethernet	125 MB/sec
Ultrium tape	320 MB/sec
PCI bus	528 MB/sec
Sun Gigaplane XB backplane	20 GB/sec

Some typical device, network, and data base rates

Wide range of types and data rates


Device Controllers

- **I/O devices have components:**
 - mechanical component
 - electronic component
- **The electronic component is the device controller**
 - may be able to handle multiple devices
- **Controller's tasks**
 - convert serial bit stream to block of bytes
 - perform error correction as necessary
 - Copy data to main memory if necessary

Memory-Mapped I/O (0)

- **Separate I/O and Memory space**
 - Address spaces for I/O and Memory are different
 - I/O port number : assigned to each control register
 - Use special I/O instructions
 - E.g. “IN REG, PORT” or “OUT PORT, REG”
- **Memory-mapped I/O**
 - Each control register is assigned a unique memory address
 - Use regular memory reference instructions to access I/O control registers
 - E.g. “store REG, Ether_Control_REG1”
- **Hybrid**
 - Both of the above schemes

Memory-Mapped I/O (1)



- **Separate I/O and memory space**
- **Memory-mapped I/O**
- **Hybrid**

Memory-Mapped I/O (0)

- **Advantages of Memory-mapped I/O**

- **No special instructions needed**

- » Device registers can be addresses in C/C++ without using assembly code

- **No special protection mechanism is needed to keep user processes from performing I/O**

- » Assign address spaces for I/O registers to protected pages

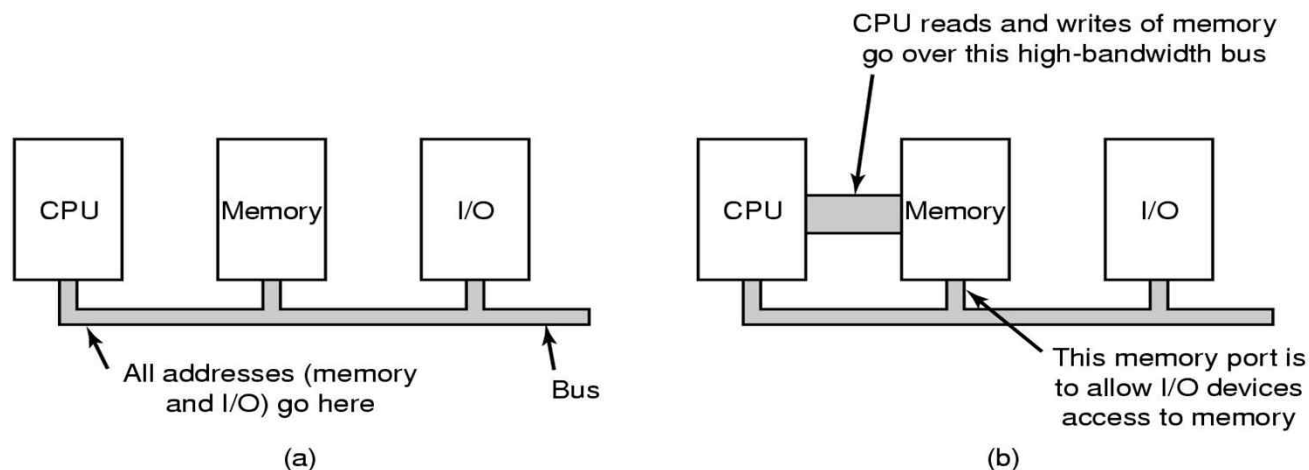
- **Every instruction that can reference memory can also reference control registers**

- » E.g. LOOP : TEST PORT_4
 BEQ READY
 BRANCH LOOP

- READY:

Memory-Mapped I/O (0)

- **Disadvantages of Memory-mapped I/O**
 - **Need hardware able to selectively disable caching**
 - » A device register should not be cached
 - **All memory modules and I/O devices must examine all memory references to see which ones to respond to**
 - » Memory bus and I/O bus are separated, we have three choices
 - CPU try memory first and if it fails, try other I/O buses
 - A snooping device on the memory bus
 - Filter addresses in the PCI bridge chip (Pentium Configuration) Fig 1-11



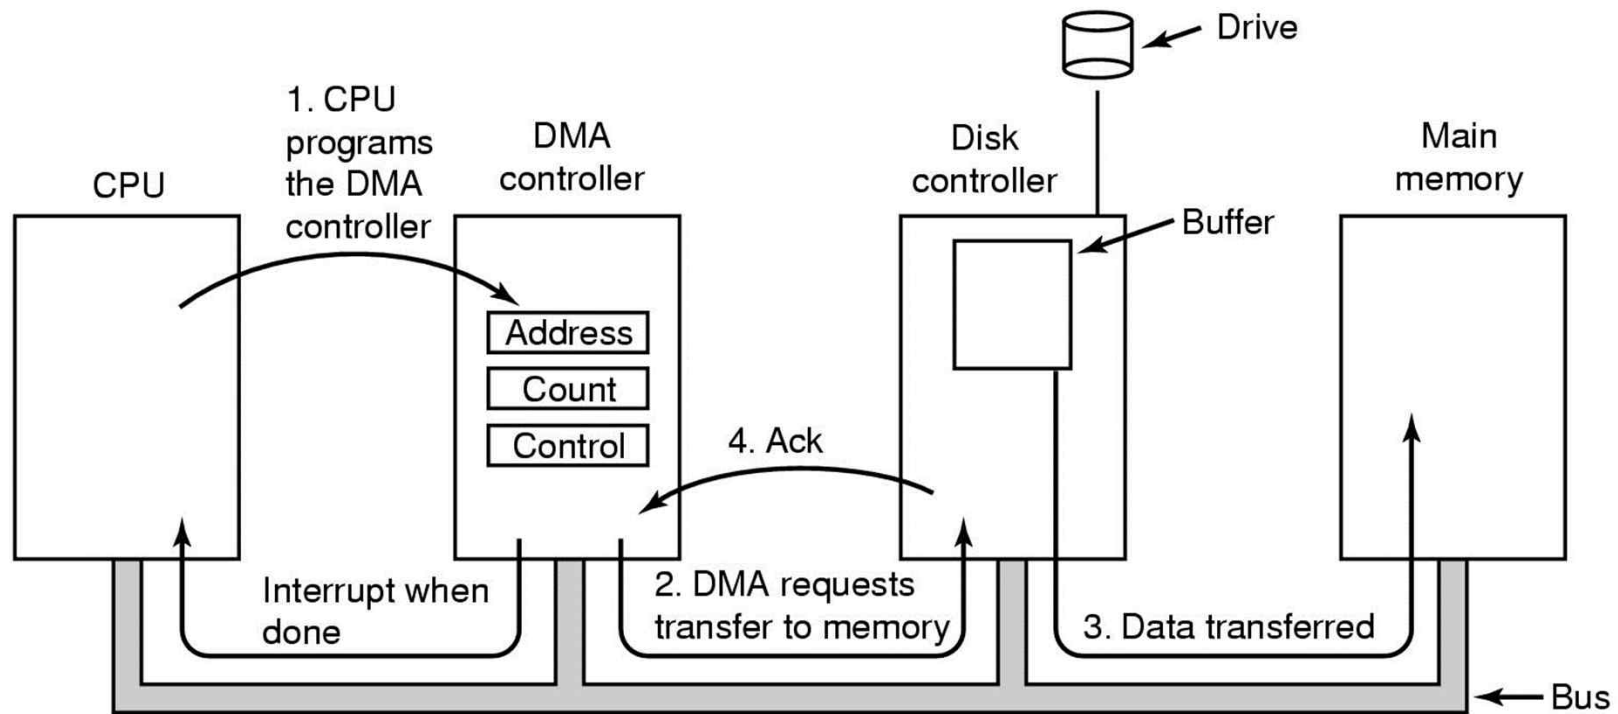
(a) A single-bus architecture

(b) A dual-bus memory architecture

Direct Memory Access (DMA)

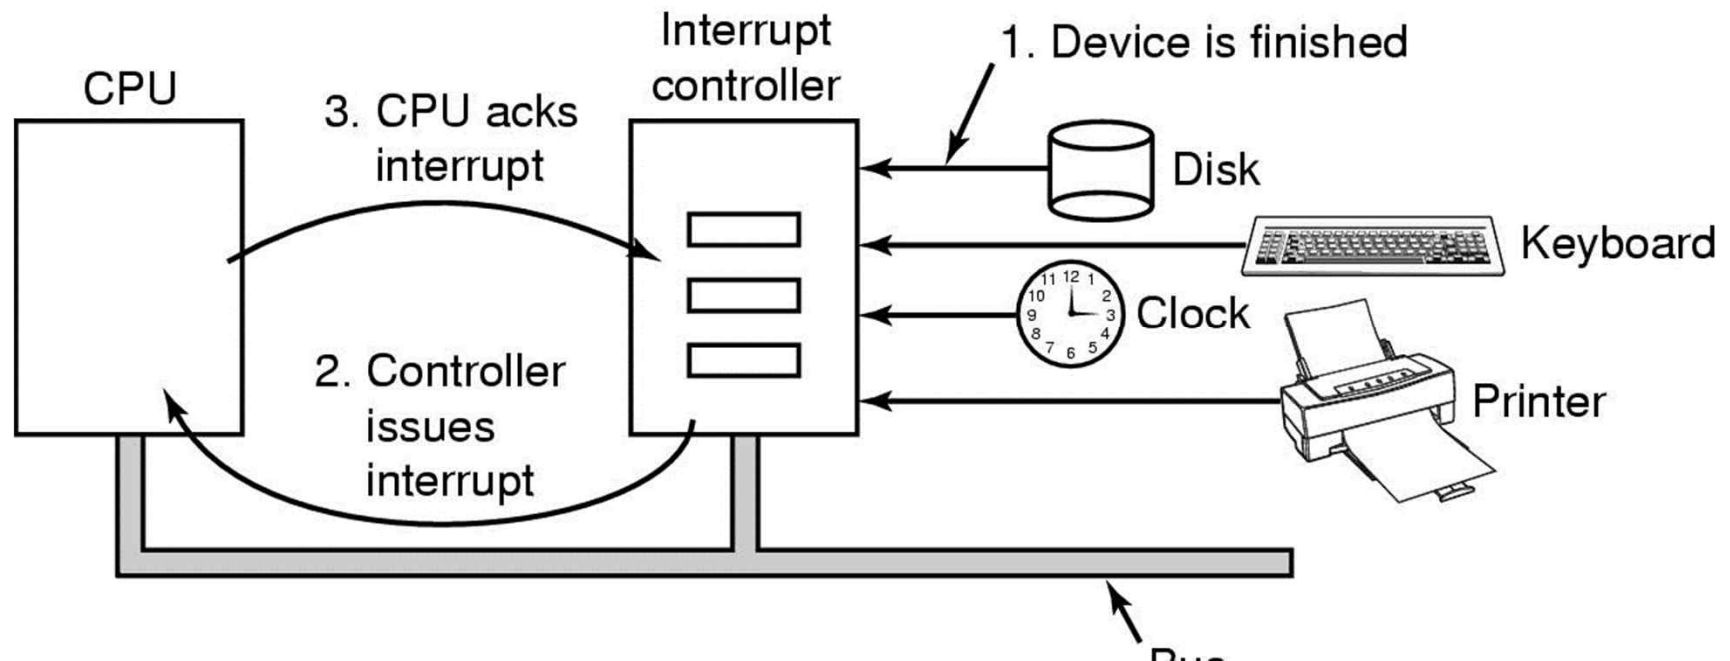
- **DMA (Direct Memory Access)**
 - **allows certain hardware devices to access memory for reading and writing independently of CPU**
 - » Without DMA, CPU typically has to be occupied for the entire time it's performing a data transfer
 - » With DMA, CPU would initiate the transfer, do other operations while the transfer is in progress, and receive an interrupt from the DMA controller once the operation has been done
- **DMA transfer mode**
 - **Cycle stealing**
 - » The controller sneaks in and steals an occasional bus cycle from the CPU to transfer a word at a time
 - **Burst mode**
 - » The controller acquire the bus and transfer multiple words
 - » It blocks the CPU or other devices during the transfer
- **Most DMA controllers use physical addresses**

Direct Memory Access (DMA)



Operation of a DMA transfer

Interrupts Revisited




How interrupts happens. Connections between devices and interrupt controller actually use interrupt lines on the bus rather than dedicated wires

Interrupts Revisited

- **Interrupt vector**
 - To fetch a new program counter for a corresponding interrupt service routine
- **Where to save hardware state information for interrupt service returns**
 - **The current user stack**
 - » Stack pointer may not even be legal
 - » May be the end of a page, causing a page fault
 - **Kernel stack**
 - » Better chance of the stack pointer being legal and in a pinned page
 - » Require change MMU contexts to switch to kernel
 - » Probably invalidate most or all of the cache and TLB
- **Precise vs. imprecise interrupts**
 - A precise interrupt that leaves the machine in a well-defined state and an imprecise interrupt does not
 - OS has to handle a large amount of internal state information generated by an imprecise interrupt

Programmed I/O (1)



Steps in printing a string

Programmed I/O (2)

```
copy_from_user(buffer, p, count);          /* p is the kernel bufer */
for (i = 0; i < count; i++) {              /* loop on every character */
    while (*printer_status_reg != READY) ;  /* loop until ready */
    *printer_data_register = p[i];          /* output one character */
}
return_to_user();
```

Writing a string to the printer using programmed I/O

- CPU do all the work - tying up the CPU full time until I/O is done
- Polling or busy waiting

Interrupt-Driven I/O

```
copy_from_user(buffer, p, count);  
enable_interrupts( );  
while (*printer_status_reg != READY) ;  
*printer_data_register = p[0];  
scheduler( );
```

(a) Code executed when print system call is made

```
if (count == 0) {  
    unblock_user( );  
} else {  
    *printer_data_register = p[i];  
    count = count - 1;  
    i = i + 1;  
}  
acknowledge_interrupt( );  
return_from_interrupt( );
```

(b) Interrupt service procedure

- **Writing a string to the printer using interrupt-driven I/O**
 - CPU does something else while waiting for I/O to become ready
 - I/O interrupts when it completes the operation and is ready
 - Interrupts take time, so frequent interrupts waste CPU time

I/O Using DMA

```
copy_from_user(buffer, p, count);  
set_up_DMA_controller( );  
scheduler( );
```

(a) Code executed when print
system call is made

```
acknowledge_interrupt( );  
unblock_user( );  
return_from_interrupt( );
```

(b) Interrupt service procedure

- **Printing a string using DMA**

- DMA controller feed the characters to the printer one at a time
- Without the CPU being bothered
- Reduces the number of interrupts from one per char to one per buffer printed

Interrupt Handlers (1)

- **Interrupt handlers are best hidden**
 - have driver starting an I/O operation block until interrupt notifies of completion
- **Interrupt procedure does its task**
 - then unblocks driver that started it
- **Steps must be performed in software after interrupt completed**
 1. **Save regs not already saved by interrupt hardware**
 2. **Set up context for interrupt service procedure**
 - Set up TLB, MMU, and a page table

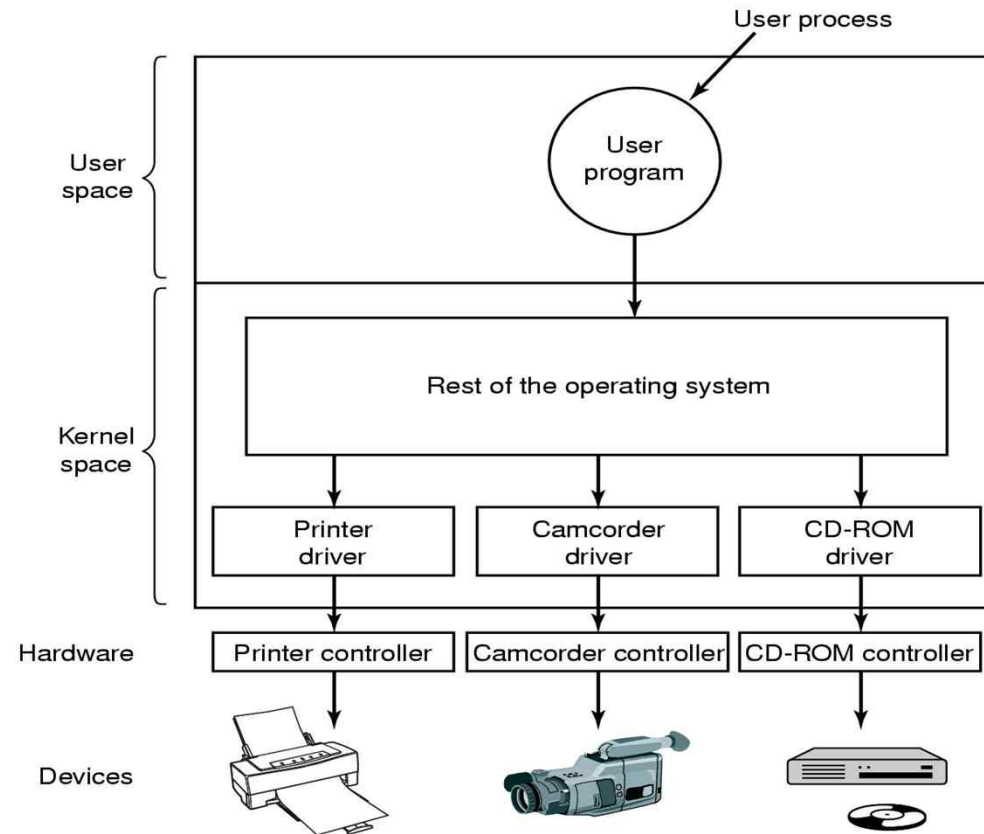
Interrupt Handlers (2)

3. **Set up stack for interrupt service procedure**
4. **Ack interrupt controller, reenable interrupts**
5. **Copy registers from where saved**
6. **Run service procedure**
7. **Set up MMU context for process to run next**
8. **Load new process' registers**
9. **Start running the new process**

Device Drivers

- Each I/O device attached to a computer needs some device-specific code for controlling the device
 - The code is called *device driver*
 - Device manufacture commonly supply drivers for OS
- Device driver normally has to be part of the OS
 - Need to access device's registers
- Device drivers are dynamically loaded to the system
 - Old approach : every driver need to be compiled into the OS
- Device drivers
 - Provide read/write funcs to upper level sw
 - Initialize the device and perform power management for the device etc.
- Device drivers have to be reentrant
 - Interrupt requiring the driver to run may happen during the driver execution

Device Drivers



- Logical position of device drivers is shown here
- Communications between drivers and device controllers goes over the bus


Disks and the OS

- **Disks are messy, messy devices**
 - errors, bad blocks, missed seeks, etc.
- **Job of OS is to hide this mess from higher-level software**
 - low-level device drivers (initiate a disk read, etc.)
 - higher-level abstractions (files, databases, etc.)
- **OS may provide different levels of disk access to different clients**
 - physical disk block (surface, cylinder, sector)
 - disk logical block (disk block #)
 - file logical (filename, block or record or byte #)

Physical Disk Structure

- **Disk components**

- platters
- surfaces
- tracks
- sectors
- cylinders
- arm
- heads



Interacting with Disks

- **In the old days...**
 - **OS would have to specify cylinder #, sector #, surface #, transfer size**
 - » I.e., OS needs to know all of the disk parameters
- **Modern disks are even more complicated**
 - **not all sectors are the same size, sectors are remapped, ...**
 - **disk provides a higher-level interface, e.g. SCSI**
 - » exports data as a logical array of blocks [0 ... N]
 - » maps **logical blocks** to cylinder/surface/sector
 - » OS only needs to name logical block #, disk maps this to cylinder/surface/sector
 - » as a result, physical parameters are hidden from OS
 - both good and bad

Example disk characteristics

- **IBM Ultrastar 36XP drive**
 - form factor: 3.5"
 - capacity: 36.4 GB
 - rotation rate: 7,200 RPM (120 RPS, musical note C3)
 - platters: 10
 - surfaces: 20
 - sector size: 512-732 bytes
 - cylinders: 11,494
 - cache: 4MB
 - transfer rate: 17.9 MB/s (inner) – 28.9 MB/s (outer)
 - full seek: 14.5 ms
 - head switch: 0.3 ms

Disk Performance

- **Performance depends on a number of steps**
 - **seek**: moving the disk arm to the correct cylinder
 - » depends on how fast disk arm can move
 - seek times aren't diminishing very quickly
 - **rotation**: waiting for the sector to rotate under head
 - » depends on rotation rate of disk
 - rates are increasing, but slowly
 - **transfer**: transferring data from surface into disk controller, and from there sending it back to host
 - » depends on density of bytes on disk
 - increasing, and very quickly
- **When the OS uses the disk, it tries to minimize the cost of all of these steps**
 - particularly seeks and rotation


Disks

Disk Hardware (1)

Parameter	IBM 360-KB floppy disk	WD 18300 hard disk
Number of cylinders	40	10601
Tracks per cylinder	2	12
Sectors per track	9	281 (avg)
Sectors per disk	720	35742000
Bytes per sector	512	512
Disk capacity	360 KB	18.3 GB
Seek time (adjacent cylinders)	6 msec	0.8 msec
Seek time (average case)	77 msec	6.9 msec
Rotation time	200 msec	8.33 msec
Motor stop/start time	250 msec	20 sec
Time to transfer 1 sector	22 msec	17 μ sec

**Disk parameters for the original IBM PC floppy disk
and a Western Digital WD 18300 hard disk**

Disk Hardware (2)




- **Physical geometry of a disk with two zones**
 - More sectors on the outer zone than the inner one
- **A possible virtual geometry for this disk**
 - View presented to the OS
 - Used by the driver software, different than the physical format

RAID


(Redundant Array of Independent Disks)

- Caching, RAM disks deal with the latency issue.
- DISKS can also be used in PARALLEL
- This is the idea behind RAIDs
 - Redundant Array of Inexpensive Disks
 - Many RAID levels (Raid0-5)



An array of inexpensive disks
(Can read 8 tracks at once)

But we have an increased
reliability problem.
If any one disk fails,
all 8 are effectively useless.



A **redundant** array
of inexpensive
disks.

RAID

- **Level 0**

- Strips (*of k sectors*) are distributed in round-robin fashion
- best with large requests; the implementation is straightforward
- worst with requests that habitually ask for data one sector at a time : no parallelism, performance gain
- No redundancy: reliability is worse than a SLED(Single Large Expensive Disk)

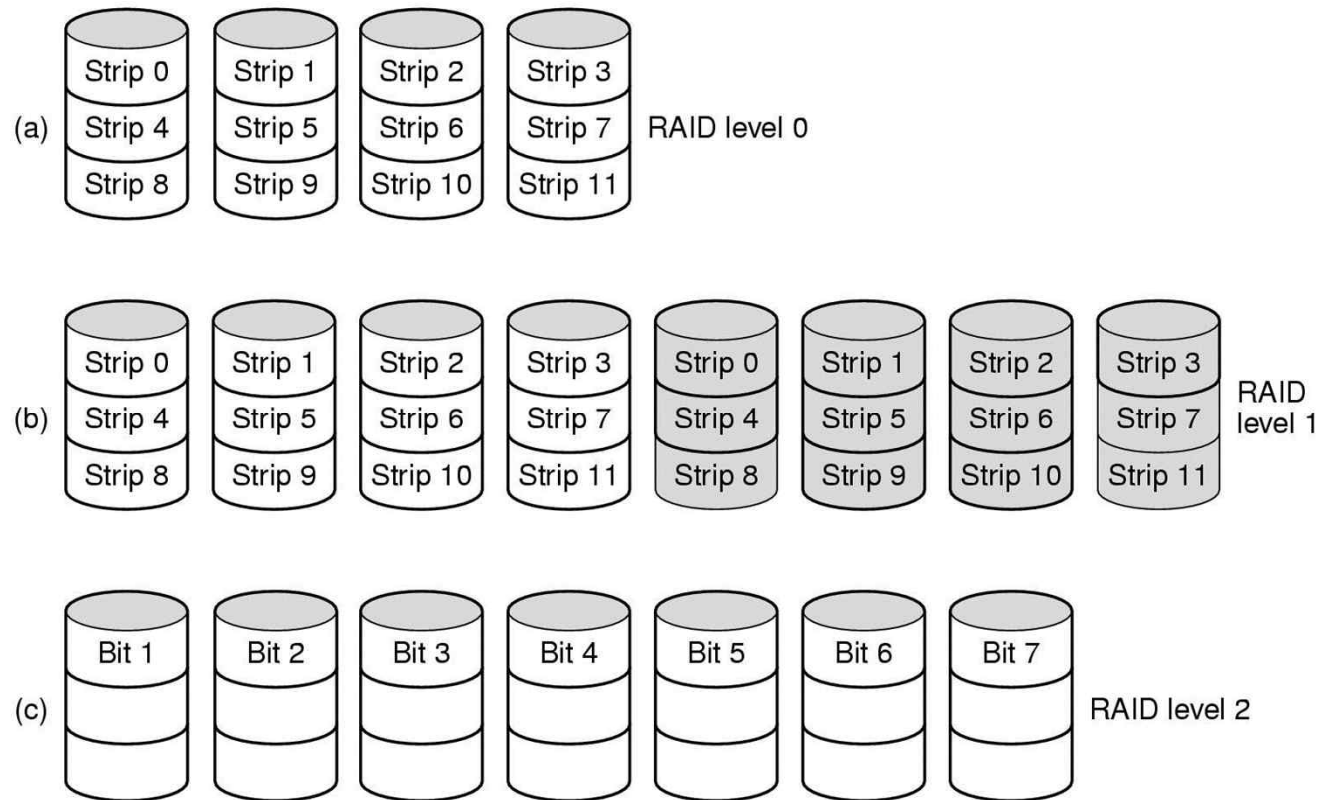
- **Level 1**

- Duplicates all the disks
- Read performance can be up to twice as good.

- **Level 2**

- Works on a word or byte basis
- E.g.. Each byte spitted into a pair of 4-bit nibbles, Hamming code added to form a 7-bit world
- Drives should be rotationally synchronized

RAID



- **Raid levels 0 through 2**
- **Backup and parity drives are shaded**

RAID

- **Level 3**

- Single parity bit is computed for each data word and written to a parity drive
- Drivers must be synchronized
- Provide full 1-bit correction by knowing the position of the bad bit (crashed driver)




- **Level 4**

- Like level 0 RAID with a strip-for-strip parity written onto an extra drive
- Performs poorly for small updates
 - » If one sector is changed, it is necessary to read all the drives in order to recalculate the parity, which must then be rewritten.
- Heavy load on the parity drive

- **Level 5**

- Eliminates bottleneck in the parity drive by distributing parity bits uniformly over all the drives, round robin fashion
- Reconstructing the contents of the failed drive is a complex process.

Disk Hardware (4)



- **Raid levels 3 through 5**
- **Backup and parity drives are shaded**

Disk Formatting

Low-level Formatting

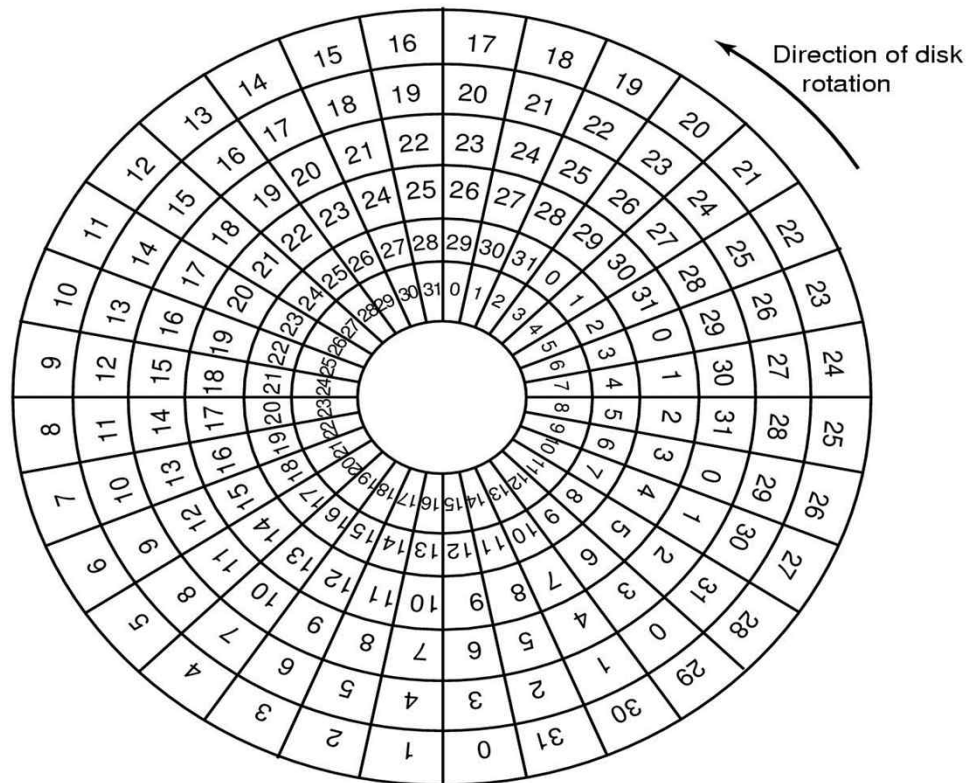


- **A disk sector**
 - **Preamble**
 - » Starts with a bit pattern that allows the hardware to recognize the start of the sector
 - » Contains the cylinder and sector numbers, etc.
 - **Data**
 - **ECC(Error-Correcting Code)**
 - » Redundant information that can be read to recover from read errors
- **Spare sectors**


Cylinder Skew

- **Cylinder skew**


- The position of sector 0 on each track is offset from the previous track when the low-level format is laid down.
- To improve performance
 - » Allows the disk to read multiple tracks in one continuous operation without losing data.




Disk Interleaving



(a)



(b)



(c)

(a) No interleaving

(b) Single interleaving

(c) Double interleaving

- To give the controller time to process the sector when consecutive sectors are read
- To avoid the need for interleaving, the controller should be able to buffer an entire track.

Sequence of Disk Formatting

- **Low-level formatting**
- **Partitioning**
 - **MBR (Master Boot Record)**
 - » Boot code
 - » Partition table
 - Starting sector and size of each partition
- **High-level formatting**
 - **Done for each partition separately**
 - **Creates a file system**
 - » Boot block, super block, free storage adm., root dir, etc.


Disk Arm Scheduling Algorithms (1)

- Time required to read or write a disk block determined by 3 factors
 1. Seek time
 2. Rotational delay
 3. Actual transfer time
- Seek time dominates
- Error checking is done by controllers

Disk Scheduling


- **Seeks are very expensive, so the OS attempts to schedule disk requests that are queued waiting for the disk**
 - **FCFS (do nothing)**
 - » reasonable when load is low
 - » long waiting time for long request queues
 - **SSTF (shortest seek time first)**
 - » minimize arm movement (seek time), maximize request rate
 - » unfairly favors middle blocks
 - **SCAN (elevator algorithm)**
 - » service requests in one direction until done, then reverse
 - » skews wait times non-uniformly (why?)
 - **C-SCAN**
 - » like scan, but only go in one direction (typewriter)
 - » uniform wait times
- **If real disk geometry != virtual geometry, what happens! - disk controller may use the algs.**

Disk Arm Scheduling Algorithms (2)



- **Shortest Seek First (SSF) disk scheduling algorithm**
 - Handles the closest request next, to minimize seek time
 - Request far from the middle may get poor service.

Disk Arm Scheduling Algorithms (3)



- **The elevator algorithm for scheduling disk requests**
 - Keeps moving in the same direction until there are no more requests in that direction, then switches direction
 - Requires one bit to keep track of the current direction: Up or Down
 - The upper bound on the total motion is fixed:
 - » $2 * (\text{number of cylinders})$

C-Scan (Circular-Scan) Algorithm


- **Similar to Scan**
- **Moves in one direction only (upward only)**
- **Smaller variance in response times than Scan**
- **The lowest-numbered cylinder is thought of as being just above the highest-numbered cylinder**

Other Optimization


- Any request to read a sector will cause that sector and much or all the rest of the current track to be read, depending upon how much space is available in the controller's cache memory
- Disk controller's cache holds blocks that have not actually been requested, while any cache maintained by the operating system will consist of blocks that were explicitly read.

Error Handling


▪ Controller : bad sector substitution



(a)



(b)



(c)

(a) A disk track with a bad sector

(b) Substituting a spare for the bad sector

(c) Shifting all the sectors to bypass the bad one

▪ OS :

- Remapping
- A secret file with all the bad sectors